

# Chronos: Concurrent Job Dispatcher in Go

Chronos is a concurrency-heavy backend system built in Go that simulates a distributed job processing system. It uses goroutines, channels, context, timeouts, retries, synchronization primitives, and graceful shutdown to manage job execution in a robust and production-grade way.

---

## Table of Contents

1. [Overview](#)
  2. [Concurrency Concepts Used](#)
  3. [Project Structure](#)
  4. [Step-by-Step Breakdown](#)
  5. Step 1: Job Definition
  6. Step 2: Job Processing
  7. Step 3: Worker Pool
  8. Step 4: Dispatcher
  9. Step 5: Graceful Shutdown
  10. Step 6: Main Integration
  11. [Final Result](#)
  12. [Next Steps](#)
- 

## Overview

Chronos simulates a backend system that:

- Accepts a list of jobs
  - Dispatches them with throttling
  - Processes them concurrently with worker goroutines
  - Applies timeout logic to each job
  - Retries failed or timed-out jobs
  - Tracks stats for monitoring
  - Handles graceful shutdown with Ctrl+C
- 

## Concurrency Concepts Used

Concept	Usage
Goroutines	For concurrent worker execution
Channels	Job queue and results

Concept	Usage
Select	Handling timeouts and rate limiting
Context	Timeout + graceful cancellation
Mutex	Safe stats updates
WaitGroup	Waiting for all workers
Atomic	Retry counters
OS Signals	Handling shutdown

---

## Project Structure

```
chronos/  
├── cmd/  
│   └── main.go  
├── internal/  
│   ├── dispatcher/  
│   │   └── dispatcher.go  
│   ├── job/  
│   │   ├── job.go  
│   │   └── processor.go  
│   ├── shutdown/  
│   │   └── shutdown.go  
│   └── worker/  
│       └── pool.go
```

---

## Step-by-Step Breakdown

### Step 1: Job Definition ( `internal/job/job.go` )

```
package job // handles job definition and logic  
  
import (  
    "fmt"           // for printing job status  
    "sync/atomic"  // for safe concurrent retry counter updates  
)  
  
type Status string // type alias for job status values
```

```

const (
    Pending Status = "pending" // job is created but not started
    Running Status = "running" // job is currently being processed
    Success Status = "success" // job completed successfully
    Failed Status = "failed" // job failed during processing
    TimedOut Status = "timed_out" // job exceeded allowed time
)

type Job struct {
    ID      int    // unique job identifier
    Payload string // content or task to be processed
    Retries int32  // number of retries
    Status  Status // current status of the job
}

// MarkRetry increments the retry count atomically
func (j *Job) MarkRetry() {
    atomic.AddInt32(&j.Retries, 1)
}

// LogStatus prints current status and retry count
func (j *Job) LogStatus() {
    fmt.Printf("[Job %d] Status: %s | Retries: %d\n", j.ID, j.Status, j.Retries)
}

```

## Step 2: Job Processing ( `internal/job/processor.go` )

```

package job // continues the job logic

import (
    "fmt"          // for logging
    "math/rand"    // to introduce random failures and durations
    "time"         // to simulate job duration
)

func Process(j *Job) error {
    fmt.Printf("[Processor] Job #%d starting with payload: %s\n", j.ID,
        j.Payload)

    // simulate processing time between 100ms to 1s
    workTime := time.Duration(rand.Intn(900)+100) * time.Millisecond
    time.Sleep(workTime)

    // simulate a 20% chance of job failure
    if rand.Float32() < 0.2 {
        return fmt.Errorf("simulated failure")
    }
}

```

```

    }

    fmt.Printf("[Processor] Job #%d completed successfully\n", j.ID)
    return nil
}

```

### Step 3: Worker Pool ( `internal/worker/pool.go` )

```

package worker // manages concurrent workers

import (
    "chronos/internal/job"
    "context"
    "fmt"
    "sync"
    "time"
)

// Stats tracks processed and failed jobs safely
type Stats struct {
    Processed int // number of successful jobs
    Failed    int // number of failed jobs
    Mutex     sync.Mutex // protects the stats from race conditions
}

func StartWorker(id int, jobs <-chan *job.Job, results chan<- *job.Job, wg
*sync.WaitGroup, stats *Stats) {
    defer wg.Done() // signal when worker is done

    for j := range jobs { // keep picking jobs until channel is closed
        j.Status = job.Running
        ctx, cancel := context.WithTimeout(context.Background(), 1*time.Second)
        resultChan := make(chan error, 1)

        // run the job processing in a goroutine to allow timeout
        go func() {
            resultChan <- job.Process(j)
        }()

        select {
        case <-ctx.Done(): // job took too long
            j.Status = job.TimedOut
            j.MarkRetry()
            stats.Mutex.Lock()
            stats.Failed++
            stats.Mutex.Unlock()

```

```

        fmt.Printf("[Worker %d] Job #%d timed out\n", id, j.ID)

    case err := <-resultChan:
        if err != nil {
            j.Status = job.Failed
            j.MarkRetry()
            stats.Mutex.Lock()
            stats.Failed++
            stats.Mutex.Unlock()
            fmt.Printf("[Worker %d] Job #%d failed: %v\n", id, j.ID, err)
        } else {
            j.Status = job.Success
            stats.Mutex.Lock()
            stats.Processed++
            stats.Mutex.Unlock()
            results <- j
        }
    }
    cancel()
}
}

```

#### Step 4: Dispatcher ( `internal/dispatcher/dispatcher.go` )

```

package dispatcher // handles job dispatch logic

import (
    "chronos/internal/job"
    "context"
    "fmt"
    "time"
)

func StartDispatcher(ctx context.Context, jobs []*job.Job, queue chan *job.Job,
rate <-chan time.Time) {
    for _, j := range jobs {
        select {
            case <-ctx.Done(): // stop dispatching if shutdown is triggered
                fmt.Println("[Dispatcher] Shutdown signal received. Stopping
dispatch...")
                close(queue)
                return
            case <-rate: // wait for rate limiter tick
                fmt.Printf("[Dispatcher] Dispatching Job #%d\n", j.ID)
                queue <- j
        }
    }
}

```

```

    }
    close(queue) // signal to workers: no more jobs
}

```

### Step 5: Graceful Shutdown ( `internal/shutdown/shutdown.go` )

```

package shutdown // handles OS interrupt signals

import (
    "context"
    "os"
    "os/signal"
    "syscall"
)

func Listen(cancel context.CancelFunc) {
    c := make(chan os.Signal, 1) // channel to capture signals
    signal.Notify(c, os.Interrupt, syscall.SIGTERM) // listen to interrupt/
    terminate

    go func() {
        <-c // wait until a signal is received
        println("\n[Shutdown] Caught interrupt. Gracefully stopping...")
        cancel() // cancel the context
    }()
}

```

### Step 6: Main Integration ( `cmd/main.go` )

```

package main

import (
    "chronos/internal/dispatcher"
    "chronos/internal/job"
    "chronos/internal/shutdown"
    "chronos/internal/worker"
    "context"
    "fmt"
    "sync"
    "time"
)

func main() {
    ctx, cancel := context.WithCancel(context.Background()) // create a global
    cancel context

```

```

defer cancel()

shutdown.Listen(cancel) // setup graceful shutdown on Ctrl+C

jobQueue := make(chan *job.Job, 10)    // main job queue
results := make(chan *job.Job, 10)     // results channel for successful
jobs
rate := time.Tick(300 * time.Millisecond) // throttle rate: 1 job every
300ms
stats := &worker.Stats{} // shared stats

// Create sample jobs
var jobList []*job.Job
for i := 1; i <= 15; i++ {
    jobList = append(jobList, &job.Job{
        ID:      i,
        Payload: fmt.Sprintf("Payload-%d", i),
        Status:  job.Pending,
    })
}

// Start the dispatcher
go dispatcher.StartDispatcher(ctx, jobList, jobQueue, rate)

// Launch worker pool
const numWorkers = 4
var wg sync.WaitGroup
for i := 1; i <= numWorkers; i++ {
    wg.Add(1)
    go worker.StartWorker(i, jobQueue, results, &wg, stats)
}

wg.Wait() // wait for all workers to complete
close(results) // no more results to receive

// Print stats
fmt.Println("\n===== FINAL REPORT =====")
fmt.Printf("Processed Jobs : %d\n", stats.Processed)
fmt.Printf("Failed Jobs   : %d\n", stats.Failed)
fmt.Println("=====")
}

```

## Final Result

- Jobs are dispatched every 300ms

- Workers handle jobs concurrently
  - Failed and timed-out jobs are tracked
  - Graceful shutdown happens on Ctrl+C
- 

## Next Steps

Feature	Description
REST API	Accept live jobs via HTTP
Retry Backoff	Use exponential backoff for retries
Persistent Queue	Redis or DB-based retrying system
Observability	Logrus for logging, Prometheus for metrics
Unit Tests	Interfaces + mock testing
Real-World Use	Convert to a microservice with endpoints